Skillberg — Deep Research Report

Prepared for: Chief Executive Officer

Contents

1. Company at a glance

Skillberg is a Lille company building what it calls the European skills graph. The legal entity is MV PROG, a SASU registered in Lille on 22 February 2024 under the software publishing classification with €10,000 of capital, with Martin Vielvoye as president. The company was incubated at EuraTechnologies.

The origin is research rather than commerce, and this is the most distinctive thing about the company. Before Skillberg, Vielvoye conducted research in artificial intelligence and computational neuroscience on how to formally represent the structure of human knowledge — specifically, how to model the fact that one skill unlocks others. That work produced the Skill Trees: prerequisite-ordered skill graphs with self-generating algorithms that compute, for any target skill, the shortest path from a learner's current position. The company's framing of the underlying epistemological claim is worth quoting in substance: deep learning makes no sense without linear algebra, Bayesian statistics and Python programming, so modelling skills as a bag of words is a category error, and modelling them as a prerequisite tree is what makes strategic HR questions answerable.

The asset is now substantially larger than the original one. Skillberg states a graph of about 158,000 canonical skills, 20,600 occupations and 18,073 cross-standard junctions spanning eight reference systems including ESCO, ROME and O*NET, from a May 2026 snapshot with monthly updates. That is a considerable expansion from the 2024 positioning, which described roughly 20,000 sub-skills per technology occupation and framed the company as a technology and IT skills repository for recruitment, training and skills mapping.

The commercial surface has two parts: infrastructure, exposed through a REST API with OpenAPI 3.1 documentation, an MCP connector available one-click in Claude, and managed graph access; and Skillberg Intelligence, a services offer applying the graph to customer workforce questions through diagnostic, mapping, analysis and roadmap work. Public descriptions of the business model have included platform licences alongside API subscription, with the learner-facing side free.

The timeline the company publishes runs: doctoral formalisation of the prerequisite model, company creation in Lille in 2024, ingestion of ESCO, ROME and O*NET into the first canonical graph with cross-standard junctions in 2025, and the public API in 2026. Public support is documented at two points: BPI France backing for the R&D phase and recruitment, and an academic research partnership structured as a CIFRE thesis, with the company describing twenty-four months of CIFRE R&D behind the graph.

On scale there is a discrepancy worth naming rather than smoothing. Skillberg's own About page presents a seven-person team, all France-based with no outsourcing of code, graph or research work, while the administrative registry shows zero employees for 2026. A CIFRE arrangement and contractor or founder-status roles would plausibly account for part of that, but the public record does not establish the employment structure. Earlier public material described a project starting in January 2024 with a paid proof of concept from the outset and a team of four by September 2024.

Questions to ask on the day:

- What changed commercially when Skillberg moved from a skills platform toward an infrastructure and API model?
- Which is currently the strongest source of repeatable revenue — platform access, API usage, or Intelligence services?
- Is the seven-person team employed by MV PROG, or does the operating model include CIFRE, contractor or other arrangements?

2. Strategic direction

The stated mission is to make skills AI explainable, sovereign and accessible, and each of those three words is doing commercial work.

Explainability is presented as a native property of the graph rather than a bolted-on feature: scores, matches and recommendations can be traced back through the graph to underlying nodes. That is a structural argument against embedding-based competitors, and it is the kind of claim that becomes more valuable as buyers face scrutiny over automated employment decisions.

Sovereignty is the second pillar, with French hosting through Scaleway and European models including Mistral positioned as differentiators. In a market where the large incumbents are American, and in a French and European buying environment increasingly attentive to that, this is a commercial argument as much as a technical one.

Interoperability is the third: combining eight occupational and skills reference systems rather than betting on one taxonomy. The cross-standard junctions are the concrete expression of that, and they are the part of the asset hardest for a competitor to replicate quickly.

Read together, the strategy is not to use AI to improve HR. It is to own a defensible skills representation layer that can be embedded into different workforce workflows — which is why the company exposes the graph by API and MCP rather than only through an application, and why it names HR as the first vertical with health, public sector and finance as intended extensions.

Two strategic tensions follow from the public position and neither is resolved publicly. The first is sequencing across verticals: the same graph may be reusable, but customer workflows, regulation, data structures and buying processes differ materially by sector, and a seven-person company can serve one sector well or several badly. The second is the relationship between infrastructure and services. Skillberg Intelligence generates revenue and market learning, both valuable at this stage, and simultaneously pulls a small team toward bespoke work at exactly the moment the infrastructure needs to become repeatable.

Questions to ask on the day:

- Which direction is receiving the most CEO attention today — building the graph, winning customers, scaling the API, or opening new verticals?
- What has been deliberately excluded from the roadmap because it would dilute the infrastructure proposition?
- How do you distinguish a durable differentiator from a feature a competitor could reproduce in a quarter?

3. Industry forces

The macro case for the category is strong and well documented. The World Economic Forum reported in January 2025 that 39% of the key skills required on the job are expected to change by 2030 and that 63% of employers identify skills gaps as a major barrier to transformation, with AI and big data, networks and cybersecurity and technological literacy among the fastest-growing skill areas. It also found that 77% of employers expect to upskill workers in response to AI while 41% expect workforce reductions where AI automates tasks.

Static taxonomies are under pressure and every serious competitor says so. TechWolf argues that surveys and self-rated skills do not scale and promotes automated inference from work activity. Neobrain markets a dynamic multilingual framework enriched from internal and external data and mapped to ROME, O*NET and ESCO. Workday describes a maintained ontology of more than 73,000 skills connected across recruiting, learning, talent management and analytics. The competitive question has moved from who has a skills library to whose representation stays current, connects to actual work, and produces decisions customers trust.

That last clause is where the category's paradox sits. Buyers want automation and productivity while demanding explainability, fairness, accountability and evidence behind recommendations — and public scrutiny of automated hiring outcomes, including litigation reported in August 2026, keeps the risk visible. This is favourable ground for a company whose core claim is traceability, provided the traceability is documented rather than asserted.

Enterprise data fragmentation is the structural problem underneath all of it. Skills intelligence in practice means reconciling employee data, job descriptions, learning content, recruiting records, organisational structures and external labour-market signals on an ongoing basis, which is why every scaled competitor emphasises integration with existing HR environments. A graph only creates value if customers can use it inside the systems they already run.

The economics are the hard part of the picture. TechWolf has raised around $55 million across three rounds and employs roughly 130 people; Neobrain has reported a €20 million Series A; Lightcast expanded through three acquisitions in 2025 spanning data, consulting and workforce transformation. A technically strong seven-person company competes against better-funded data, integration and distribution footprints, which means the path from technical differentiation to repeatable revenue has to be unusually clear rather than merely plausible.

Questions to ask on the day:

- Which pressure is creating the strongest buying urgency among your customers today?
- Where does the category still fail to translate into measurable business outcomes?
- Which is hardest to compete against — data depth, distribution, integrations, trust, or customer education?

4. Regulatory and compliance frame

The regulatory picture moved in July and August 2026, and the movement matters commercially because Skillberg positions regulatory readiness as a differentiator.

The AI Act entered into force on 1 August 2024, with prohibited practices and AI literacy applying from 2 February 2025 and governance and general-purpose AI obligations following in August 2025. Enforcement powers took effect on 2 August 2026. But the Digital Omnibus on AI, Regulation (EU) 2026/1744, in force since 27 July 2026, deferred the Annex III high-risk obligations — which include employment uses such as recruitment, selection, ranking and decisions affecting work relationships or employee management — from 2 August 2026 to 2 December 2027, with embedded Annex I systems moving to 2 August 2028. Article 50 transparency obligations applied from 2 August 2026.

There is a specific and checkable point here: Skillberg's current public wording on the regulatory timeline does not align with the Commission's current position on when Annex III employment rules apply. That is a small factual discrepancy on a page whose entire purpose is to establish regulatory credibility, which makes it worth more attention than its size suggests.

The more durable strategic reading is that the proposition is not "opaque HR AI is now illegal". It is that legal classification depends on intended purpose and actual use — the same underlying skills technology carries different obligations depending on whether it supports search and reporting, recommendations, ranking, or decisions that materially affect people — and that customers therefore need traceable, governable and auditable skills intelligence regardless of which side of December 2027 they are on.

Data protection sits alongside it and bites sooner. CNIL guidance states that recruitment processing falls within GDPR and must follow specified purpose, lawfulness and minimisation, and it addresses automated candidate sorting, ranking and evaluation directly. Article 22 safeguards can require human intervention, an opportunity to express a viewpoint and a route to contest a decision. CNIL has also stated that organisations using personal data for AI training must establish an appropriate legal basis and consider the provenance and reuse context of the data — which makes data provenance a recurring commercial issue wherever CVs, employee data, job descriptions or customer datasets enrich a graph or matching logic. Article 4 AI literacy obligations apply to providers and deployers, so compliance includes organisational practice and user understanding rather than product configuration alone.

Questions to ask on the day:

- Which regulatory question do prospects raise most often?
- Do you define compliance at the graph level, the product level, or the use-case level?
- What evidence do you personally need before making a public claim about compliance, sovereignty or explainability?

5. The chief executive's world at Skillberg

CEO research consistently identifies six dimensions of the role — setting strategy, aligning the organisation, leading the top team, working with the board, representing the company externally, and managing one's own time and energy — alongside resource allocation, capital allocation, organisation design and talent decisions. Calendar research finds the job dominated by meetings and interactions across internal and external constituencies, with attention as the scarce resource.

At seven people that framework compresses. The chief executive personally carries more commercial, product, talent and external-relations work than the research describes, and here there is an additional feature: the founder is also the inventor of the core intellectual property and the named author of the research behind it. That makes him the technical credibility in every room as well as the commercial lead, which is an asset in a category where buyers are sceptical and a constraint on everything that depends on his time.

The likely measurement categories at this stage are customer traction, growth, cash efficiency, product progress, talent and strategic credibility, though no public source sets out an actual scorecard.

The recurring friction in this role is not a shortage of information. It is the work of making information decision-ready — and for a skills-intelligence company the inputs are unusually diverse. Customer feedback, labour-market shifts, competitor moves, taxonomy changes, regulatory developments, product performance, pipeline evidence and internal capability can all bear on the same decision, and each of them changes on its own schedule. The eight reference systems in the graph are themselves eight external sources that move independently.

Questions to ask on the day:

- Which recurring decision consumes disproportionate time because the underlying facts are scattered?
- What do you most often have to ask someone to find, verify or reformat personally?
- Which executive work feels repetitive even though the decision itself is strategic?

6. Peer signals

TechWolf is the closest direct comparison and the most instructive on the competitive structure. It positions itself as a data layer for the AI era of work, combining skills, work and market intelligence and embedding that intelligence into systems including SAP and Workday. It has raised roughly $55 million across three rounds and employs around 130 people from Ghent. The sharpest fact is who invested: its $42.75 million Series B in 2024 included Workday Ventures, ServiceNow Ventures and SAP alongside financial investors. The platforms an independent skills-infrastructure company would most need to integrate with are already shareholders in a competitor — which is a distribution problem before it is a product problem.

Neobrain positions around skills, performance and career growth with a dynamic framework enriched by internal data and external market insight, and has reported a €20 million Series A. The signal is that customers are being offered a complete operating model around skills — framework, mapping, assessment, development and mobility — rather than a taxonomy.

Lightcast acquired Rhetorik, Skill Collective and Simply during 2025, positioning the Skill Collective deal as combining skills data with consulting capability to take customers from strategy to execution. Scaled players are trying to own the whole path from data to transformation, which is the same instinct behind Skillberg Intelligence, at a very different scale.

Workday Skills Cloud is embedded in the HCM platform with an ontology of more than 73,000 skills connected across recruiting, learning, talent management and analytics. The lesson is that skills intelligence gains strategic value where HR decisions are already made, and that standalone intelligence carries a heavier integration and adoption burden.

Eightfold launched an AI Recruiter Agent and AI Interviewer in 2025 while publicly emphasising AI governance, compliance and human-led AI, including ISO/IEC 42001-related certification. Trust and auditability now travel alongside automation rather than being sold separately — which is Skillberg's argument, made by a much larger company with certification to back it.

Questions to ask on the day:

- Which competitor is genuinely difficult to displace rather than merely adjacent?
- Where does the moat sit — graph structure, data, research, interoperability, regulatory positioning, or customer relationships?
- SAP, Workday and ServiceNow are investors in TechWolf. Does that close doors, or does it validate the category?

7. Friction map

The evidence base is external and it moves. Eight reference systems, a labour market that reprices skills continuously, competitors updating their propositions, and a regulatory timeline that changed twice in 2026 — the external surface a skills-infrastructure company must track is larger than any single team's data, and none of it changes on the company's schedule.

Claims need evidence, not assertion. Explainability, sovereignty and interoperability are the product, which means each is a statement a customer, regulator or journalist may test. The company's own regulatory wording currently sitting slightly out of step with the Commission's timeline is a small instance of a structural problem: public claims decay unless someone owns keeping them current.

The same truth gets framed differently for different audiences. Customers, product, commercial teams, partners and investors describe the same capability in different terms and measures, and reconciling those framings is unglamorous work that falls to whoever wrote the original.

Maintenance is the product, not a chore attached to it. Monthly graph updates across eight standards is an ongoing obligation, and the competitive claim of currency means it cannot be deprioritised when something more urgent arrives.

Services pull against infrastructure. The Intelligence offer generates revenue and market learning while consuming the scarce engineering and research attention that makes the API repeatable, and at seven people those are the same people.

Competitive comparison has become multi-dimensional. Products, data assets, distribution, capital, integration footprint and positioning all matter simultaneously, and a feature-by-feature view no longer describes the contest.

Questions to ask on the day:

- Which of these is a weekly occurrence rather than an occasional annoyance?
- Where do you have to chase three or more sources before one answer can be trusted?

- Which repeated information task would you most like to stop personally supervising?

Challenge cards

Card 1 of 2

Title: Infrastructure and services, same seven people

The challenge: The company sells two things from one team. The API, MCP connector and managed graph access are infrastructure that only pays if it becomes repeatable and largely self-serve. Skillberg Intelligence is diagnostic, mapping and roadmap work that pays now, teaches the team what customers actually need, and consumes the same research and engineering attention the infrastructure requires.

Who feels it: The chief executive, who allocates the attention and is personally involved in both. The engineering and research team, whose weeks get claimed by whichever is more urgent. And customers on the infrastructure side, whose roadmap items wait behind delivery work.

Why it persists: Both are correct at this stage. Services revenue funds the graph and produces market learning no amount of desk research would, and at seven people a services engagement is genuinely the fastest route to understanding a new sector. But every hour of bespoke work is an hour not spent making the next customer need less bespoke work, and the tension does not resolve itself — it has to be decided, repeatedly.

What solved looks like: A stated rule about what Intelligence engagements are for, so each one is accepted or declined against a purpose rather than against capacity, and so what is learned in delivery reliably becomes graph or product rather than staying in the engagement.

Evidence base: Skillberg publicly presents two commercial surfaces — infrastructure through REST API, MCP for Claude and managed graph access, and Skillberg Intelligence applying the graph to customer workforce questions through diagnostic, mapping, analysis and roadmap work. The company presents a seven-person France-based team with no outsourcing of code, graph or research work. Its stated vertical ambition extends from HR into health, public sector and finance, each with different workflows, regulation and data structures. Comparable players have taken the same path at very different scale: Lightcast acquired consulting capability in 2025 specifically to move customers from skills strategy to execution.

Open question: Of the Intelligence engagements delivered so far, how many produced something that made it into the graph or the API, and how many stayed with the customer?

Card 2 of 2

Title: Keeping the claims current

The challenge: Statements about explainability, data practices, sovereignty, market position and regulatory fit have to remain accurate as the underlying evidence moves. Regulations change, competitors reposition, the graph grows, and customer diligence returns to the same topics — so the same facts get re-established repeatedly, and the public wording drifts out of step with the evidence behind it.

Who feels it: The chief executive, who makes the public claims and is the technical authority behind them. Also engineering and research, asked to confirm technical statements; commercial colleagues, whose material derives from those statements; and prospective customers' compliance functions, who test them.

Why it persists: The sources change independently and none of them announce themselves. The EU AI Act timetable moved twice in 2026 — enforcement powers from 2 August, and Annex III employment obligations deferred to 2 December 2027 under Regulation (EU) 2026/1744 — while CNIL guidance, eight reference taxonomies and competitor positioning all move on separate schedules. For a company whose entire differentiation is explainability and regulatory credibility, a stale claim is not a copywriting problem.

What solved looks like: Repeated customer, partner and public questions are answered with confidence that the wording is current and supported, and someone other than the chief executive owns the currency of the public claims. The regulatory position on the website is right because it is maintained, not because it was correct when written.

Evidence base: Skillberg's current public regulatory wording is not fully aligned with the European Commission's current timeline for Annex III employment rules. The company positions explainability, European hosting through Scaleway and European models including Mistral, and cross-standard interoperability as its three differentiators, each of which is a testable claim. The graph is stated at 158,000 canonical skills across eight reference systems from a May 2026 snapshot with monthly updates, so the headline figures themselves have a shelf life. CNIL guidance covers automated candidate sorting and evaluation, AI training data provenance, and Article 22 human-oversight safeguards. Eightfold, a much larger competitor, has pursued ISO/IEC 42001-related certification to evidence equivalent claims.

Open question: Which external claim or recurring diligence question is hardest for you personally to keep current and defensible?

Where this document is thin

- No usable revenue, recurring revenue, profitability, cash or valuation figures are public; the 2025 accounts are filed as confidential.
- No priced institutional funding round for Skillberg was identified. Public support is documented as BPI France backing for the R&D and recruitment phase and an academic CIFRE research partnership, which are not the same as a venture round.
- Beneficial ownership beyond the SASU form and Martin Vielvoye as president is not established publicly.
- The seven-person team presented publicly and the registry's zero-employee figure for 2026 are not reconciled. CIFRE and contractor arrangements would plausibly account for part of it, but the employment structure is not documented.
- No public customer list gives enough detail to establish customer concentration, renewal rates, contract sizes or geographic mix, and no public evidence establishes which vertical contributes most revenue.
- The split of business between API usage, platform licences and Intelligence services is not public.
- Graph metrics, performance figures and infrastructure claims on the company's site are company-published and have not been independently validated. The company's placement in the 2026 TalentCLEF evaluation is externally published, but a single benchmark position is not a general quality measure.
- No public evidence establishes deployment at material scale outside France.
- No public CEO scorecard or formal objectives were identified.
- Whether Skillberg's own systems fall within Annex III depends on intended purpose and customer use, which cannot be determined from public sources and probably varies by deployment.